

# Course Introduction & History of Database Systems

---

CMPT 843, SPRING 2019

JIANNAN WANG

<https://sfu-db.github.io/dbsystems/>

# Introduce Yourself

---

What's your name?

Where are you from?

M.Sc. or Ph.D.? Which year?

What do you want to get out of the course?

---

# Course Introduction

# Why this course?

---

## Objective 1. Stand on the shoulders of giants!

- A Half Century DB Research
- VLDB = Very Large Data Base (founded in 1975)
- Four Turing Awards
- \$50 billion Market Per Year

## Objective 2. Master essential skills for being a researcher

- Reading Papers
- Giving Talks
- Reviewing Papers
- Asking Questions

# Prerequisites

---

## Undergraduate Database Systems Courses

- **Testing yourself:** Relational Model, SQL, Query Optimization, Transaction, Concurrency Control, ACID, etc.

If not, you have to spend extra time on

- Textbook: “[Database Management Systems](#)”
- Online Courses: [Stanford](#), [Berkeley](#)

---

**Part 1: Traditional Database Systems and Techniques** (20 papers)

**Part 2: Modern Database Systems and Techniques** (18 papers)

**39  
papers**

# Part 1: Traditional Database Systems and Techniques (before 2000)

---

## Background

1. Database Systems: Achievements and Opportunities (1990)
2. The Asilomar Report on Database Research (1998)

# Part 1: Traditional Database Systems and Techniques (before 2000)

---

## Data Model

- 3. A Relational Model of Data for Large Shared Data Banks (1970)
- 4. What Goes Around Comes Around (1960-1970, Sec I~IV only)

# Part 1: Traditional Database Systems and Techniques (before 2000)

---

## Traditional DBMS

5. A History and Evaluation of System R (1981)
6. The Design of Postgres (1986)
7. The Gamma database machine project (1990)
8. An Overview of Data Warehousing and OLAP Technology (1997)

# Part 1: Traditional Database Systems and Techniques (before 2000)

---

## Transaction Management

9. [Granularity of Locks and Degrees of Consistency in a Shared Data Base \(1976, Part 1 only\)](#)
10. [Granularity of Locks and Degrees of Consistency in a Shared Data Base \(1976, Part 2 only\)](#)
11. [On Optimistic Methods for Concurrency Control \(1981\)](#)
12. [Concurrency control performance modeling: alternatives and implications \(1987\)](#)

# Part 1: Traditional Database Systems and Techniques (before 2000)

---

## Query Optimization

13. Access Path Selection in a Relational Database Management System (1979)
14. The Volcano Optimizer Generator: Extensibility and Efficient Search (1993)
15. Efficient mid-query re-optimization of sub-optimal query execution plans (1998)
16. Eddies: Continuously Adaptive Query Processing (2000)

# Part 1: Traditional Database Systems and Techniques (before 2000)

---

## Interactive Analytics

17. [Data Cube: A Relational Aggregation Operator Generalizing Group-by, Cross-tab, and Sub-totals \(1997\)](#)
18. [An Array-Based Algorithm for Simultaneous Multidimensional Aggregates \(1997\)](#)
19. [New Sampling-Based Summary Statistics for Improving Approximate Query Answers \(1998\)](#)
20. [Informix under CONTROL: Online Query Processing \(2000\)](#)

# Part 2: Modern Database Systems and Techniques (after 2000)

---

## Background

21. Challenges and Opportunities with Big Data (2011)

# Part 2: Modern Database Systems and Techniques (after 2000)

---

## MapReduce and Beyond

- 22. [MapReduce: Simplified Data Processing on Large Clusters \(2003\)](#)
- 23. [A Comparison of Approaches to Large-Scale Data Analysis \(2009\)](#)
- 24. [Resilient Distributed Datasets: A Fault-tolerant Abstraction for In-memory Cluster Computing \(2012\)](#)
- 25. [Spark SQL: Relational Data Processing in Spark \(2015\)](#)

# Part 2: Modern Database Systems and Techniques (after 2000)

---

## Column Store

26. [C-store: A Column-oriented DBMS \(2005\)](#)
27. [Dremel: Interactive Analysis Of Web-Scale Datasets \(2010\)](#)
28. [Column-Stores vs. Row-Stores: How Different Are They Really?\(2012\)](#)

# Part 2: Modern Database Systems and Techniques (after 2000)

---

## NoSQL

29. [Chord: A Scalable Peer-to-peer Lookup Protocol for Internet Applications \(2003\)](#)
30. [Bigtable: A Distributed Storage System for Structured Data \(2006\)](#)
31. [Dynamo: Amazon's Highly Available Key-Value Store \(2007\)](#)
32. [CAP Twelve Years Later: How the "Rules" Have Changed \(2012\)](#)

# Part 2: Modern Database Systems and Techniques (after 2000)

---

## NewSQL

33. [OLTP Through the Looking Glass, and What We Found There \(2008\)](#)
  34. [Hekaton: SQL Server's Memory-optimized OLTP Engine \(2013\)](#)
  35. [Efficiently Compiling Efficient Query Plans for Modern Hardware \(2011\)](#)
  36. [Scalable SQL and NoSQL data stores \(2010\)](#)
-

# Part 2: Modern Database Systems and Techniques (after 2000)

---

## ML and SQL

37. Accelerating Machine Learning Inference with Probabilistic Predicates (2018)
38. NoScope: Optimizing Deep CNN-Based Queries over Video Streams at Scale (2017)
39. Learning to Optimize Join Queries With Deep Reinforcement Learning (2018)
40. The Case for Learned Index Structures (2018)

# References

---

## Papers to read

- [Redbook - 5th Edition](#) (Peter Bailis, Joseph M. Hellerstein, Michael Stonebraker)
- [Readings in Database](#) (Reynold Xin)

## Related Graduate Courses

- [CS286: Implementation of Database Systems](#) (UC Berkeley, Fall 2014)
- [EECS 584: Advanced Database Management Systems](#) (UMichigan, 2015 Fall)
- [Big Data Systems](#) (Columbia, 2016 Spring)
- [15-799: Advanced Topics in Database Systems](#) (CMU, 2013 Fall)

# Skills

---

Reading Papers

Giving Talks

Reviewing Papers

Asking Questions

# How you will be trained

---

## Reading Papers

- 40 Papers

## Giving Talks

- 2 Paper (**20min+10 min Q&A**)

## Reviewing Papers

- 2 Papers

## Asking Questions

- Asking at least 10 questions in the Q&A sessions

# Grading

---

Paper Presentation: 20%

Paper Review: 20%

Participation: 10%

Questions: 10%

Blog Post: 10%

Final Project: 30% (2% plan + 14% poster + 14% report)

# What's next

---

Fill in the form by the end of Monday 1/7  
<https://goo.gl/g3y735>

---

# History of Database Systems

# Database Systems in a Half Century

## (1960s – 2010s)

When	What
Early 1960 – Early 1970	The Navigational Database Empire
Mid 1970 – Mid 1980	The Database World War I
Mid 1980 – Early 2000	The Relational Database Empire
Mid 2000 – Now	The Database World War II

### References.

- <https://en.wikipedia.org/wiki/Database#History>
- [What Goes Around Comes Around \(Michael Stonebraker, Joe Hellerstein\)](#)
- [40 Years VLDB Panel](#)

# The Navigational Database Empire

## (Early 1960 – Early 1970)

### Data Model

1. How to organize data
2. How to access data

### Navigational Data Model

1. Organize data into a multi-dimensional space (i.e., A space of records)
2. Access data by following pointers between records

### Inventor: Charles Bachman

1. The 1973 ACM Turing Award
2. Turing Lecture: "The Programmer As Navigator"



# The Navigational Database Empire

## (Early 1960 – Early 1970)

---

### Representative Navigational Database Systems

- Integrated Data Store (IDS), 1964, GE
- Information Management System (IMS), 1966, IBM
- Integrated Database Management System (IDMS), 1973, Goodrich

### CODASYL

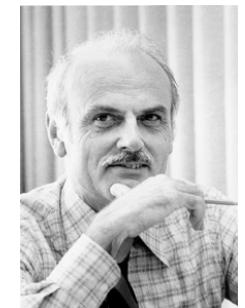
- Short for “Conference/Committee on Data Systems Languages”
- Define navigational data model as standard database interface (1969)

# The Birth of Relational Model

---

## Ted Codd

- Born in 1923
- PHD in 1965
- "A Relational Model of Data for Large Shared Data Banks" in 1970



## Relational Model

- Organize data into a collection of relations
- Access data by a declarative language (i.e., tell me what you want, not how to find it)

Data  
Independence

# The Database World War I: Background

---

## One Slide (Navigational Model)

- Led by Charles Bachman (1973 ACM Turing Award)
- Has built mature systems
- Dominated the database market

## The other Slide (Relational Model)

- Led by Ted Codd (mathematical programmer, IBM)
- A theoretical paper with no system built
- Little support from IBM

# The Database World War I: Three Big Campaigns

---

1. Which data model is better in **theory**?  
(Mid 1970)
  
2. Which data model is better in **practice**?  
(Late 1970 – Early 1980)
  
3. Which data model is better in **business**?  
(Early 1980 – Mid 1980)

# The “Theory” Campaign

---

A **debate** at ACM SIGFIDET (precursor of SIGMOD)  
1974

Navigational model is bad

- **Data Organization:** So complex
- **Data Access:** No declarative language

Relational model is bad

- **Data Organization:** A special case of navigational model
- **Data Access:** No system proof that declarative language is viable

# The “Practice” Campaign

---

## The Big Question

- Can a relational database system perform as good as a navigational system?

## System prototypes

- Ingres at UC Berkeley (early and pioneering) 
- System R at IBM (arguably got more stuff “right”) 

## The System R Team

- Query Optimization (Patricia P. Griffiths et al.)
- SQL ( Donald D. Chamberlin et al.)
- Transaction (Jim Gray et al.)

# The “Business” Campaign

---

Commercialization of Relational Database Systems

Not as easy as we thought

Three reasons (required) that led relational database systems to win

- The minicomputer revolution (1977)
- Competing products (e.g. IDMS) could not be ported to the minicomputer
- Relational front end was not added to navigational database systems

# What Can We Learn?

---

**Lesson 1.**

The winning of theory  $\neq$  The winning of practice

**Lesson 2.**

The winning of practice  $\neq$  The winning of business

**Lesson 3.**

Everyone can get a chance to win

# The Relational Database Empire (Mid1980 – Early 2000)

---

## Parallel and distributed DBs (1980 – 1990)

- SystemR\*, Distributed Ingres, Gamma, etc.

## Object-oriented DB (1980 – 1990)

- Objects: Data/Code Integration
- Extensibility: User-defined functions, User-defined data types

## MySQL and PostgreSQL (1990s)

- Widely used open-source relational DB systems

# The Database World War II: Background

---

## Internet Boom (Early 2000)

- Larger data volume that cannot be fit in a single machine
- Faster data updates that cannot be handled by a single machine

Commercial distributed database systems are expensive 😞

Open-source database systems do not support distributed computing well 😞

# The Database World War II: Two Big Campaigns

---

1. Which is better for large-scale data analysis:  
MapReduce vs. Traditional Database Systems?
2. Which is better for distributed transaction processing:  
NoSQL vs. Traditional Database Systems?

Let's find the answers through this course! 😊